

An AI real-estate agent that says how sure it is — before it says what to buy.



Most AI chat tools answer immediately, whether or not they actually have enough information to. Home Compass is built around the opposite bet: don't guess, don't invent, and show your work — before you say anything to someone about where they should live.

01 THE PROBLEM

Two things go wrong with AI real-estate chat today.

First, confidence. A chatbot that's only heard "I'm looking for a house" will still cheerfully suggest neighborhoods — staying silent isn't in its nature, it was built to always have an answer. That's fine for trivia. It's a problem when the "answer" is where someone should live.

Second, trust. Ask it about a neighborhood's crime rate or school quality and it can state a specific number that sounds authoritative and simply isn't real — a well-documented failure mode called hallucination. In real estate that isn't a quirk, it's exposure: characterizing a neighborhood touches fair housing law, whether the words came from a licensed agent or a machine.

The two most common complaints about AI advice — "it made that up" and "it didn't actually listen to

me" — both trace back to the same root cause: the system was never designed to say "I'm not sure yet."

02 THE FIX

Two rules the agent isn't allowed to break.

1 Don't answer until you're sure.

Before it searches anything, the agent tracks how well it actually understands the request — budget, location, must-haves, timeline — and shows that as a live reading, the way a fuel gauge shows a fraction rather than "full" or "empty." Below a set threshold, it asks one more targeted question instead of guessing.

2 Don't say anything you can't back up — or aren't allowed to say.

Every factual claim — price, crime count, income data, market trend — is required to trace back to a named source and date. Separately, every finished answer is reviewed against a rulebook modeled on fair housing regulations before it's shown, the way a compliance officer would review a human agent's language — not asked to self-police after the fact.

03 THE DIFFERENCE, CONCRETELY


What changes in the actual conversation.

SITUATION	TYPICAL AI CHAT	HOME COMPASS
You've only said "I need a place to live"	✗ Guesses anyway, offers neighborhoods on the spot	✓ Asks one targeted follow-up first

You ask "is this area safe?"	✗ Answers in vague, confident-sounding terms	✓ Cites the government source and year, or says the data isn't available
Before a neighborhood is recommended	✗ No compliance review of any kind	✓ Checked against a fair-housing rulebook first
"How sure is this thing, right now?"	✗ No way to tell	✓ Shown as a live, visible reading throughout

04 HOW A CONVERSATION ACTUALLY RUNS

Five steps, the same shape every time.



05 WHY IT'S A BUSINESS, NOT JUST A FEATURE

Where the value actually lands.

A human agent's time is the scarcest resource in this business, and most of it goes to buyers who aren't ready yet. An assistant that only hands off a lead

once it's genuinely confident it understands that buyer is doing real qualifying work, not just deflecting questions. And because the compliance check runs on every single answer — not just the ones a human remembers to be careful about — it's a meaningfully lower-risk way to put AI in front of a home buyer than a general-purpose chatbot with real estate data bolted on.

PATH A

Embedded assistant

Licensed to brokerages and listing portals as a white-labeled chat experience on their own site.

PATH B

Warm handoff

The assistant qualifies the buyer to a confidence threshold, then routes a warmer lead to a human agent than a raw contact form ever produces.

PATH C

Platform feature

Bundled into an existing real-estate SaaS product as the trust-and-compliance differentiator versus competitors' bolt-on chatbots.

06 WHERE THIS STANDS TODAY

What's real, and what a launch would still need.

BUILT & WORKING

The full mechanism above runs end to end today: the confidence gate, the sourced search and ranking, and the compliance review — tested against a public sample listings dataset and a realistic rulebook built to mirror real fair-housing rules for testing purposes.

STILL NEEDED

A live, licensed listings feed in place of the sample dataset; the real rulebook replaced with one reviewed and signed off by actual counsel; and the standard hardening any prototype goes through before it handles real users — none of it an open research problem, all of it ordinary next-stage work.

*"Don't guess. Don't invent. Show your work."
— the whole product, in one sentence.*